

```
In [1]: #Setting up GPU for training
import os
os.environ["CUDA_DEVICE_ORDER"]="PCI_BUS_ID"
os.environ["CUDA_VISIBLE_DEVICES"]="0" #model will be trained on GPU 0 NVIDIA RTX 3060
```

```
In [2]: import keras
from matplotlib import pyplot as plt
import numpy as np
import gzip
%matplotlib inline
from keras.models import Model
from keras.optimizers import RMSprop
from keras.layers import Input, Dense, Flatten, Dropout, Reshape, Conv2D, MaxPooling2D, UpSampling2D, Concatenate, GlobalAveragePooling2D
from keras.layers import add, concatenate, merging
from keras.models import Model, Sequential
from keras.callbacks import ModelCheckpoint
from keras.optimizers import Adadelta, RMSprop, SGD, Adam
from keras import regularizers
from keras import backend as K
from keras.utils import to_categorical
```

Extracting the MNIST Fashion dataset

1. Extractin the dataset.
2. reading into buffer
3. convertinf to tuple of (no_of_imgs, widht,height, channel) (60k,28,28)

```
In [3]: def extract_data(filename, num_images):
    with gzip.open(filename) as bytestream:
        bytestream.read(16)
        buf = bytestream.read(28 * 28 * num_images)
        data = np.frombuffer(buf, dtype=np.uint8).astype(np.float32)
        data = data.reshape(num_images, 28,28)
    return data
```

```
In [4]: train_data = extract_data('train-images-idx3-ubyte.gz', 60000)
test_data = extract_data('t10k-images-idx3-ubyte.gz', 10000)
```

```
In [5]: def extract_labels(filename, num_images):
    with gzip.open(filename) as bytestream:
        bytestream.read(8)
        buf = bytestream.read(1 * num_images)
        labels = np.frombuffer(buf, dtype=np.uint8).astype(np.int64)
    return labels
```

```
In [6]: train_labels = extract_labels('train-labels-idx1-ubyte.gz', 60000)
test_labels = extract_labels('t10k-labels-idx1-ubyte.gz', 10000)
```

```
In [7]: print("Training set (images) shape: {shape}".format(shape=train_data.shape))
print("Test set (images) shape: {shape}".format(shape=test_data.shape))
```

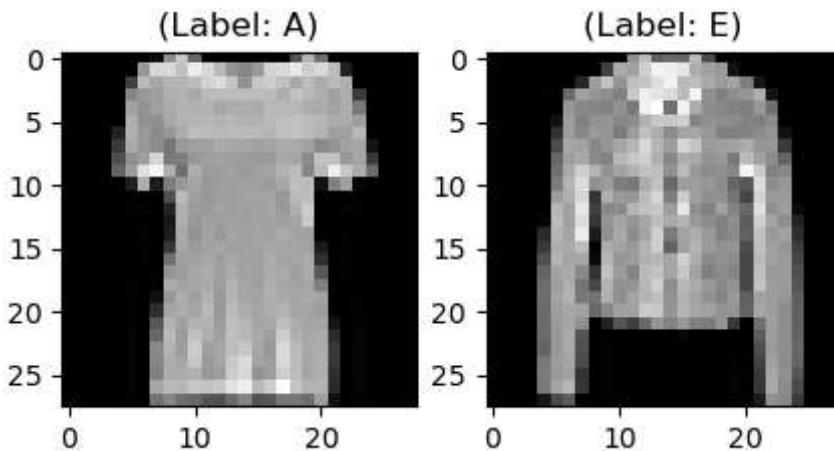
Training set (images) shape: (60000, 28, 28)
Test set (images) shape: (10000, 28, 28)

```
In [9]: # Create dictionary of target classes
label_dict = {
    0: 'A',
    1: 'B',
    2: 'C',
    3: 'D',
    4: 'E',
    5: 'F',
    6: 'G',
    7: 'H',
    8: 'I',
    9: 'J',
}
```

```
In [10]: plt.figure(figsize=[5,5])
plt.subplot(121)
curr_img = np.reshape(train_data[10], (28,28))
curr_lbl = train_labels[10]
plt.imshow(curr_img, cmap='gray')
plt.title("(Label: " + str(label_dict[curr_lbl]) + ")")

plt.subplot(122)
curr_img = np.reshape(test_data[10], (28,28))
curr_lbl = test_labels[10]
plt.imshow(curr_img, cmap='gray')
plt.title("(Label: " + str(label_dict[curr_lbl]) + ")")
```

Out[10]: Text(0.5, 1.0, '(Label: E)')



```
In [11]: # reshaping the dataset with no_of_imgs, width, height, channel
train_data = train_data.reshape(-1, 28,28, 1)
test_data = test_data.reshape(-1, 28,28, 1)
train_data.shape, test_data.shape
```

Out[11]: ((60000, 28, 28, 1), (10000, 28, 28, 1))

```
In [12]: # normalizing the dataset
train_data.dtype, test_data.dtype
```

Out[12]: (dtype('float32'), dtype('float32'))

```
In [13]: np.max(train_data), np.max(test_data)
```

Out[13]: (255.0, 255.0)

```
In [14]: train_data = train_data / np.max(train_data)
test_data = test_data / np.max(test_data)
```

```
In [15]: np.max(train_data), np.max(test_data)
```

Out[15]: (1.0, 1.0)

```
In [16]: #splitting the dataset for train validation
from sklearn.model_selection import train_test_split
train_X,valid_X,train_ground,valid_ground = train_test_split(train_data,
                                                               train_data,
                                                               test_size=0.2,
                                                               random_state=13)
```

```
In [21]: batch_size = 64
epochs = 5
inChannel = 1
x, y = 28, 28
input_img = Input(shape = (x, y, inChannel))
num_classes = 10
```

```
In [18]: def encoder(input_img):
    #input = 28 x 28 x 1 (wide and thin)
    conv1 = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img) #28 x 28
    conv1 = BatchNormalization()(conv1)
    conv1 = Conv2D(32, (3, 3), activation='relu', padding='same')(conv1)
    conv1 = BatchNormalization()(conv1)
    pool1 = MaxPooling2D(pool_size=(2, 2))(conv1) #14 x 14 x 32
    conv2 = Conv2D(64, (3, 3), activation='relu', padding='same')(pool1) #14 x 14 x 64
    conv2 = BatchNormalization()(conv2)
    conv2 = Conv2D(64, (3, 3), activation='relu', padding='same')(conv2)
    conv2 = BatchNormalization()(conv2)
    pool2 = MaxPooling2D(pool_size=(2, 2))(conv2) #7 x 7 x 64
    conv3 = Conv2D(128, (3, 3), activation='relu', padding='same')(pool2) #7 x 7 x 128
    conv3 = BatchNormalization()(conv3)
    conv3 = Conv2D(128, (3, 3), activation='relu', padding='same')(conv3)
    conv3 = BatchNormalization()(conv3)
    conv4 = Conv2D(256, (3, 3), activation='relu', padding='same')(conv3) #7 x 7 x 256
    conv4 = BatchNormalization()(conv4)
    conv4 = Conv2D(256, (3, 3), activation='relu', padding='same')(conv4)
    conv4 = BatchNormalization()(conv4)
    return conv4

def decoder(conv4):
    #decoder
    conv5 = Conv2D(128, (3, 3), activation='relu', padding='same')(conv4) #7 x 7 x 128
    conv5 = BatchNormalization()(conv5)
    conv5 = Conv2D(128, (3, 3), activation='relu', padding='same')(conv5)
    conv5 = BatchNormalization()(conv5)
    conv6 = Conv2D(64, (3, 3), activation='relu', padding='same')(conv5) #7 x 7 x 64
    conv6 = BatchNormalization()(conv6)
    conv6 = Conv2D(64, (3, 3), activation='relu', padding='same')(conv6)
    conv6 = BatchNormalization()(conv6)
    up1 = UpSampling2D((2,2))(conv6) #14 x 14 x 64
    conv7 = Conv2D(32, (3, 3), activation='relu', padding='same')(up1) # 14 x 14 x 32
```

```
conv7 = BatchNormalization()(conv7)
conv7 = Conv2D(32, (3, 3), activation='relu', padding='same')(conv7)
conv7 = BatchNormalization()(conv7)
up2 = UpSampling2D((2,2))(conv7) # 28 x 28 x 32
decoded = Conv2D(1, (3, 3), activation='sigmoid', padding='same')(up2) # 28 x 28 x 1
return decoded
```

```
In [19]: autoencoder = Model(input_img, decoder(encoder(input_img)))
autoencoder.compile(loss='mean_squared_error', optimizer = RMSprop())
```

```
In [20]: autoencoder.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d (Conv2D)	(None, 28, 28, 32)	320
batch_normalization (BatchNormalization)	(None, 28, 28, 32)	128
conv2d_1 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_1 (BatchNormalization)	(None, 28, 28, 32)	128
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_2 (BatchNormalization)	(None, 14, 14, 64)	256
conv2d_3 (Conv2D)	(None, 14, 14, 64)	36928
batch_normalization_3 (BatchNormalization)	(None, 14, 14, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_4 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_4 (BatchNormalization)	(None, 7, 7, 128)	512
conv2d_5 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_5 (BatchNormalization)	(None, 7, 7, 128)	512
conv2d_6 (Conv2D)	(None, 7, 7, 256)	295168
batch_normalization_6 (BatchNormalization)	(None, 7, 7, 256)	1024
conv2d_7 (Conv2D)	(None, 7, 7, 256)	590080
batch_normalization_7 (BatchNormalization)	(None, 7, 7, 256)	1024
conv2d_8 (Conv2D)	(None, 7, 7, 128)	295040
batch_normalization_8 (BatchNormalization)	(None, 7, 7, 128)	512
conv2d_9 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_9 (BatchNormalization)	(None, 7, 7, 128)	512

```

hNormalization)

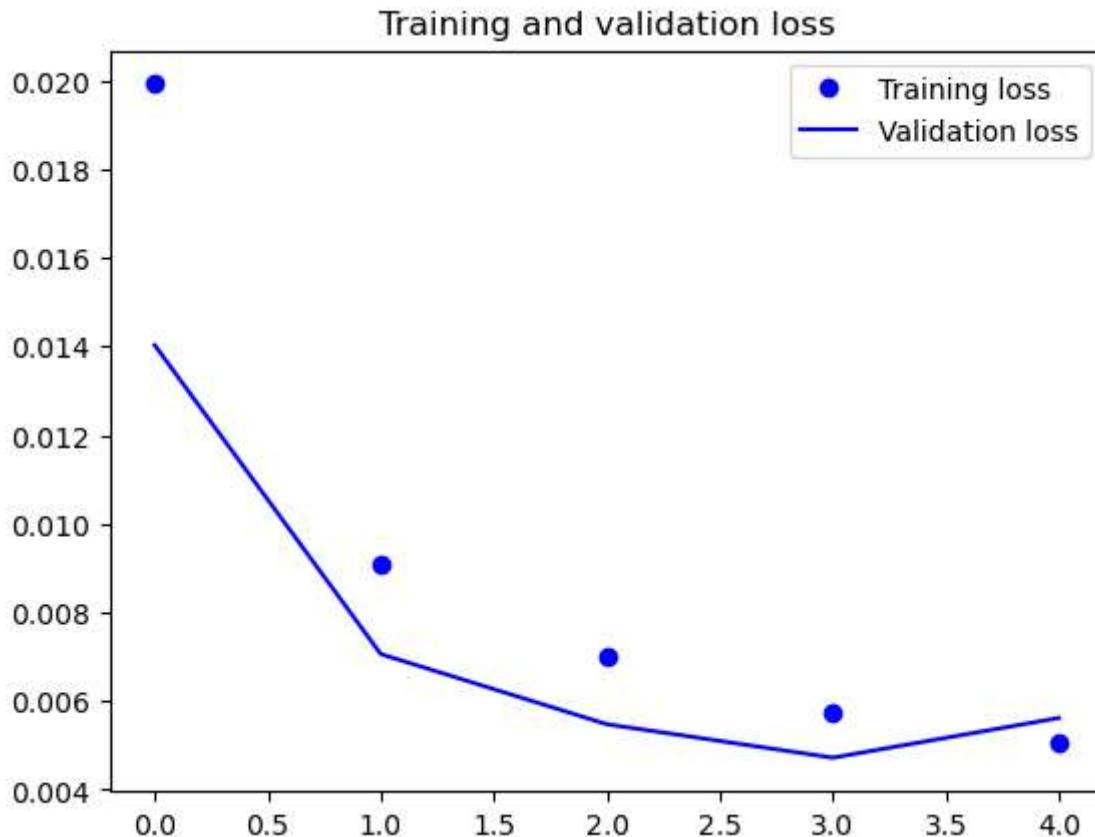
conv2d_10 (Conv2D)           (None, 7, 7, 64)      73792
batch_normalization_10 (BatchNormalization)        (None, 7, 7, 64)      256
conv2d_11 (Conv2D)           (None, 7, 7, 64)      36928
batch_normalization_11 (BatchNormalization)        (None, 7, 7, 64)      256
up_sampling2d (UpSampling2D) (None, 14, 14, 64)     0
conv2d_12 (Conv2D)           (None, 14, 14, 32)    18464
batch_normalization_12 (BatchNormalization)        (None, 14, 14, 32)    128
conv2d_13 (Conv2D)           (None, 14, 14, 32)    9248
batch_normalization_13 (BatchNormalization)        (None, 14, 14, 32)    128
up_sampling2d_1 (UpSampling2D) (None, 28, 28, 32)   0
conv2d_14 (Conv2D)           (None, 28, 28, 1)     289
=====
Total params: 1,758,657
Trainable params: 1,755,841
Non-trainable params: 2,816

```

```
In [22]: autoencoder_train = autoencoder.fit(train_X, train_ground, batch_size=batch_size, epochs=5)
Epoch 1/5
750/750 [=====] - 414s 546ms/step - loss: 0.0199 - val_loss: 0.0140
Epoch 2/5
750/750 [=====] - 236s 315ms/step - loss: 0.0091 - val_loss: 0.0071
Epoch 3/5
750/750 [=====] - 234s 313ms/step - loss: 0.0070 - val_loss: 0.0055
Epoch 4/5
750/750 [=====] - 236s 314ms/step - loss: 0.0057 - val_loss: 0.0047
Epoch 5/5
750/750 [=====] - 234s 313ms/step - loss: 0.0050 - val_loss: 0.0056
```

```
In [24]: loss = autoencoder_train.history['loss']
val_loss = autoencoder_train.history['val_loss']
epochs = range(5)
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
```

```
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



In [25]: `autoencoder.save_weights('autoencoder.h5')`

In [26]: `# Change the Labels from categorical to one-hot encoding
train_Y_one_hot = to_categorical(train_labels)
test_Y_one_hot = to_categorical(test_labels)`

`print('Original label:', train_labels[0])
print('After conversion to one-hot:', train_Y_one_hot[0])`

Original label: 9

After conversion to one-hot: [0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]

In [27]: `train_X,valid_X,train_label,valid_label = train_test_split(train_data,train_Y_one_hot,`

In [28]: `train_X.shape,valid_X.shape,train_label.shape,valid_label.shape`

Out[28]: `((48000, 28, 28, 1), (12000, 28, 28, 1), (48000, 10), (12000, 10))`

In [29]: `# creating the classification model with same as encoder network and assigning the same weights`
`def encoder(input_img):`
 `#encoder`
 `#input = 28 x 28 x 1 (wide and thin)`
 `conv1 = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img) #28 x 28`
 `conv1 = BatchNormalization()(conv1)`
 `conv1 = Conv2D(32, (3, 3), activation='relu', padding='same')(conv1)`
 `conv1 = BatchNormalization()(conv1)`
 `pool1 = MaxPooling2D(pool_size=(2, 2))(conv1) #14 x 14 x 32`

```

conv2 = Conv2D(64, (3, 3), activation='relu', padding='same')(pool1) #14 x 14 x 64
conv2 = BatchNormalization()(conv2)
conv2 = Conv2D(64, (3, 3), activation='relu', padding='same')(conv2)
conv2 = BatchNormalization()(conv2)
pool2 = MaxPooling2D(pool_size=(2, 2))(conv2) #7 x 7 x 64
conv3 = Conv2D(128, (3, 3), activation='relu', padding='same')(pool2) #7 x 7 x 128
conv3 = BatchNormalization()(conv3)
conv3 = Conv2D(128, (3, 3), activation='relu', padding='same')(conv3)
conv3 = BatchNormalization()(conv3)
conv4 = Conv2D(256, (3, 3), activation='relu', padding='same')(conv3) #7 x 7 x 256
conv4 = BatchNormalization()(conv4)
conv4 = Conv2D(256, (3, 3), activation='relu', padding='same')(conv4)
conv4 = BatchNormalization()(conv4)
return conv4

```

In [30]:

```

# adding the FCNN as an head to encoder and decoder network
# using functional API for layer sequencing
def fc(enco):
    flat = Flatten()(enco)
    den = Dense(128, activation='relu')(flat)
    out = Dense(num_classes, activation='softmax')(den)
    return out

```

In [31]:

```

encode = encoder(input_img)
full_model = Model(input_img,fc(encode))

```

In [32]:

```

for l1,l2 in zip(full_model.layers[:19],autoencoder.layers[0:19]):
    l1.set_weights(l2.get_weights())

```

In [33]:

```

autoencoder.get_weights()[0][1]

```

```
Out[33]: array([[[ -1.54724228e-04,   8.53857622e-02,   1.40017807e-01,
   3.70581564e-03,  -1.22385100e-01,  -9.03333724e-02,
  -1.42092910e-02,   9.53786969e-02,   5.75400293e-02,
  9.30630192e-02,   8.25757384e-02,  -8.72964738e-05,
  -3.02403383e-02,  -6.17197976e-02,   1.08013727e-01,
  -6.76150322e-02,   3.28186192e-02,   4.71103564e-02,
  -9.78118256e-02,   1.47139803e-01,   1.22941025e-01,
  1.51436487e-02,  -9.05194283e-02,   3.97238359e-02,
  -1.15449227e-01,   6.40309453e-02,   8.10088292e-02,
  -3.40074711e-02,  -1.34132281e-01,  -1.19687431e-01,
  -8.79123956e-02,   1.08853234e-02]],

[[ 1.91540793e-01,   1.69836655e-01,   5.17360261e-03,
  9.11975279e-02,  -7.66993240e-02,  -9.51103047e-02,
  -4.71934676e-02,   2.36582085e-01,  -3.66107076e-02,
  1.46808177e-01,   1.19001545e-01,   1.35637134e-01,
  1.28245085e-01,  -3.67860310e-02,  -3.72047164e-02,
  6.57114238e-02,  -2.11761147e-02,  -1.54801562e-01,
  -8.48870128e-02,  -5.92605136e-02,   3.50146443e-02,
  4.84344959e-02,  -2.60472689e-02,   1.74833670e-01,
  -4.50035511e-03,   1.51497841e-01,   5.65965995e-02,
  -1.27874196e-01,   1.68987781e-01,   9.66395438e-02,
  1.07529864e-01,   2.21004054e-01]],

[[ 3.11180577e-02,   1.43947929e-01,   1.48395285e-01,
  -5.66556975e-02,   1.27784222e-01,  -1.94676250e-01,
  -1.12775555e-02,   1.49947613e-01,   7.85046890e-02,
  4.91623543e-02,  -1.16249183e-02,   1.75571308e-01,
  1.36704937e-01,   1.91662177e-01,  -6.67674765e-02,
  1.46571219e-01,  -1.78442031e-01,  -5.34193926e-02,
  5.86946197e-02,  -1.33571669e-01,   3.41310240e-02,
  -1.57855704e-01,   3.69145907e-02,   7.27996230e-02,
  -5.03489859e-02,   2.08528087e-01,  -2.12112620e-01,
  1.44237895e-02,  -1.03738010e-01,   4.29570861e-02,
  -9.62668881e-02,  -6.54554814e-02]]], dtype=float32)
```

```
In [34]: full_model.get_weights()[0][1]
```

```
Out[34]: array([[[[-1.54724228e-04,  8.53857622e-02,  1.40017807e-01,
   3.70581564e-03, -1.22385100e-01, -9.03333724e-02,
  -1.42092910e-02,  9.53786969e-02,  5.75400293e-02,
  9.30630192e-02,  8.25757384e-02, -8.72964738e-05,
  -3.02403383e-02, -6.17197976e-02,  1.08013727e-01,
  -6.76150322e-02,  3.28186192e-02,  4.71103564e-02,
  -9.78118256e-02,  1.47139803e-01,  1.22941025e-01,
  1.51436487e-02, -9.05194283e-02,  3.97238359e-02,
  -1.15449227e-01,  6.40309453e-02,  8.10088292e-02,
  -3.40074711e-02, -1.34132281e-01, -1.19687431e-01,
  -8.79123956e-02,  1.08853234e-02]],

[[ 1.91540793e-01,  1.69836655e-01,  5.17360261e-03,
  9.11975279e-02, -7.66993240e-02, -9.51103047e-02,
  -4.71934676e-02,  2.36582085e-01, -3.66107076e-02,
  1.46808177e-01,  1.19001545e-01,  1.35637134e-01,
  1.28245085e-01, -3.67860310e-02, -3.72047164e-02,
  6.57114238e-02, -2.11761147e-02, -1.54801562e-01,
  -8.48870128e-02, -5.92605136e-02,  3.50146443e-02,
  4.84344959e-02, -2.60472689e-02,  1.74833670e-01,
  -4.50035511e-03,  1.51497841e-01,  5.65965995e-02,
  -1.27874196e-01,  1.68987781e-01,  9.66395438e-02,
  1.07529864e-01,  2.21004054e-01]],

[[ 3.11180577e-02,  1.43947929e-01,  1.48395285e-01,
  -5.66556975e-02,  1.27784222e-01, -1.94676250e-01,
  -1.12775555e-02,  1.49947613e-01,  7.85046890e-02,
  4.91623543e-02, -1.16249183e-02,  1.75571308e-01,
  1.36704937e-01,  1.91662177e-01, -6.67674765e-02,
  1.46571219e-01, -1.78442031e-01, -5.34193926e-02,
  5.86946197e-02, -1.33571669e-01,  3.41310240e-02,
  -1.57855704e-01,  3.69145907e-02,  7.27996230e-02,
  -5.03489859e-02,  2.08528087e-01, -2.12112620e-01,
  1.44237895e-02, -1.03738010e-01,  4.29570861e-02,
  -9.62668881e-02, -6.54554814e-02]]], dtype=float32)
```

```
In [35]: for layer in full_model.layers[0:19]:
    layer.trainable = False
```

```
In [36]: full_model.compile(loss=keras.losses.categorical_crossentropy, optimizer=keras.optimiz
```

```
In [37]: full_model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
<hr/>		
input_2 (InputLayer)	[(None, 28, 28, 1)]	0
conv2d_15 (Conv2D)	(None, 28, 28, 32)	320
batch_normalization_14 (BatchNormalization)	(None, 28, 28, 32)	128
conv2d_16 (Conv2D)	(None, 28, 28, 32)	9248
batch_normalization_15 (BatchNormalization)	(None, 28, 28, 32)	128
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_17 (Conv2D)	(None, 14, 14, 64)	18496
batch_normalization_16 (BatchNormalization)	(None, 14, 14, 64)	256
conv2d_18 (Conv2D)	(None, 14, 14, 64)	36928
batch_normalization_17 (BatchNormalization)	(None, 14, 14, 64)	256
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 64)	0
conv2d_19 (Conv2D)	(None, 7, 7, 128)	73856
batch_normalization_18 (BatchNormalization)	(None, 7, 7, 128)	512
conv2d_20 (Conv2D)	(None, 7, 7, 128)	147584
batch_normalization_19 (BatchNormalization)	(None, 7, 7, 128)	512
conv2d_21 (Conv2D)	(None, 7, 7, 256)	295168
batch_normalization_20 (BatchNormalization)	(None, 7, 7, 256)	1024
conv2d_22 (Conv2D)	(None, 7, 7, 256)	590080
batch_normalization_21 (BatchNormalization)	(None, 7, 7, 256)	1024
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 128)	1605760
dense_1 (Dense)	(None, 10)	1290
<hr/>		
Total params: 2,782,570		

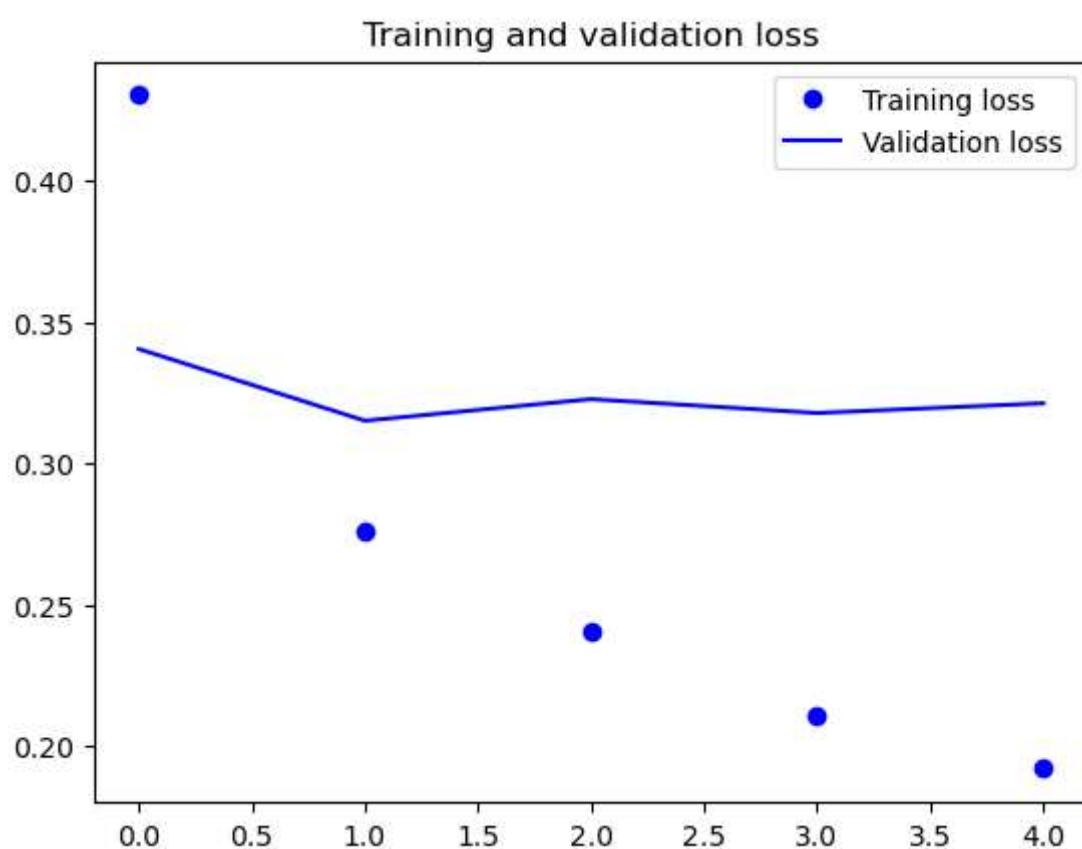
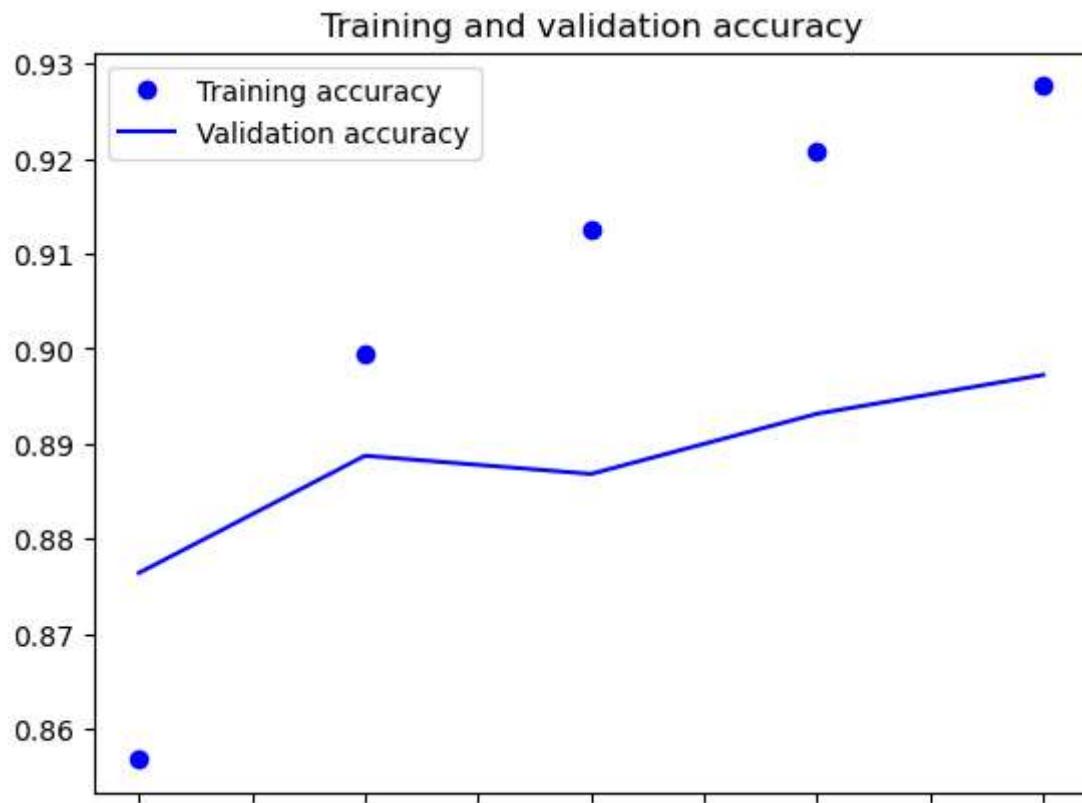
```
Trainable params: 1,607,050
Non-trainable params: 1,175,520
```

```
In [38]: classify_train = full_model.fit(train_X, train_label, batch_size=64, epochs=5, verbose=1)

Epoch 1/5
750/750 [=====] - 68s 90ms/step - loss: 0.4303 - accuracy: 0.8568 - val_loss: 0.3404 - val_accuracy: 0.8764
Epoch 2/5
750/750 [=====] - 65s 87ms/step - loss: 0.2759 - accuracy: 0.8995 - val_loss: 0.3151 - val_accuracy: 0.8888
Epoch 3/5
750/750 [=====] - 66s 88ms/step - loss: 0.2401 - accuracy: 0.9126 - val_loss: 0.3227 - val_accuracy: 0.8868
Epoch 4/5
750/750 [=====] - 66s 88ms/step - loss: 0.2105 - accuracy: 0.9207 - val_loss: 0.3178 - val_accuracy: 0.8932
Epoch 5/5
750/750 [=====] - 69s 91ms/step - loss: 0.1919 - accuracy: 0.9276 - val_loss: 0.3212 - val_accuracy: 0.8972
```

```
In [39]: full_model.save_weights('autoencoder_classification.h5')
```

```
In [42]: accuracy = classify_train.history['accuracy']
val_accuracy = classify_train.history['val_accuracy']
loss = classify_train.history['loss']
val_loss = classify_train.history['val_loss']
epochs = range(len(accuracy))
plt.plot(epochs, accuracy, 'bo', label='Training accuracy')
plt.plot(epochs, val_accuracy, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



```
In [43]: test_eval = full_model.evaluate(test_data, test_Y_one_hot, verbose=0)
```

```
In [44]: print('Test loss:', test_eval[0])
print('Test accuracy:', test_eval[1])
```

```
Test loss: 0.34199246764183044
Test accuracy: 0.8942999839782715
```

```
In [45]: predicted_classes = full_model.predict(test_data)
313/313 [=====] - 15s 48ms/step
```

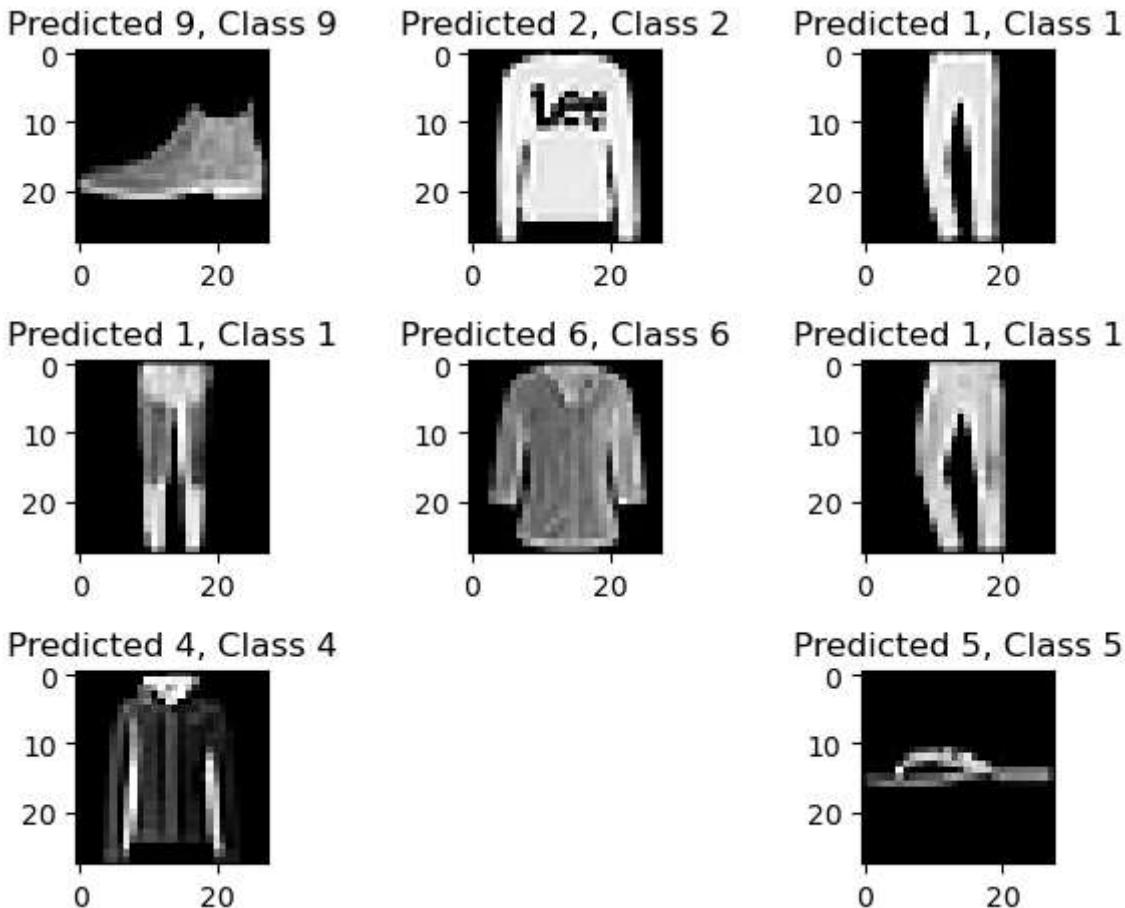
```
In [53]: predicted_classes = np.argmax(np.round(predicted_classes), axis=1)
```

```
In [54]: predicted_classes.shape, test_labels.shape
```

```
Out[54]: ((10000,), (10000,))
```

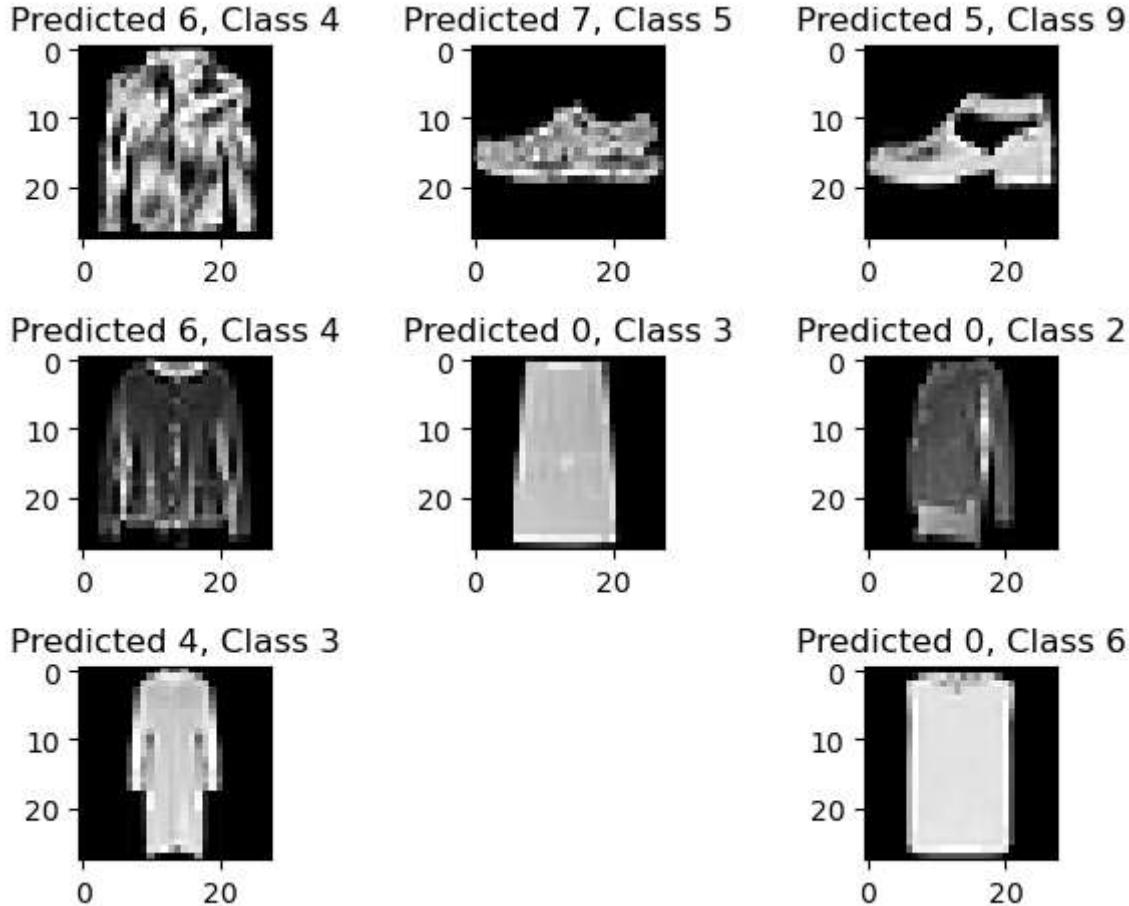
```
In [55]: correct = np.where(predicted_classes==test_labels)[0]
print("Found %d correct labels" % len(correct))
for i, correct in enumerate(correct[:9]):
    plt.subplot(3,3,i+1)
    plt.imshow(test_data[correct].reshape(28,28), cmap='gray', interpolation='none')
    plt.title("Predicted {}, Class {}".format(predicted_classes[correct], test_labels[correct]))
    plt.tight_layout()
```

Found 8921 correct labels



```
In [56]: incorrect = np.where(predicted_classes!=test_labels)[0]
print("Found %d incorrect labels" % len(incorrect))
for i, incorrect in enumerate(incorrect[:9]):
    plt.subplot(3,3,i+1)
    plt.imshow(test_data[incorrect].reshape(28,28), cmap='gray', interpolation='none')
    plt.title("Predicted {}, Class {}".format(predicted_classes[incorrect], test_labels[incorrect]))
    plt.tight_layout()
```

Found 1079 incorrect labels



```
In [58]: from sklearn.metrics import classification_report
target_names = ["Class {}".format(i) for i in range(num_classes)]
print(classification_report(test_labels, predicted_classes, target_names=target_names))
```

	precision	recall	f1-score	support
Class 0	0.74	0.89	0.81	1000
Class 1	0.99	0.97	0.98	1000
Class 2	0.81	0.85	0.83	1000
Class 3	0.90	0.90	0.90	1000
Class 4	0.83	0.82	0.82	1000
Class 5	0.99	0.96	0.97	1000
Class 6	0.82	0.61	0.70	1000
Class 7	0.94	0.97	0.95	1000
Class 8	0.97	0.98	0.97	1000
Class 9	0.97	0.96	0.97	1000
accuracy			0.89	10000
macro avg	0.89	0.89	0.89	10000
weighted avg	0.89	0.89	0.89	10000

In []: